

# DATA STRUCTURES AND ALGORITHMS

Name : Hrudhay Reddy Garisa

## Table of Contents

Question 1 .....	2
Answer to Question 1.....	2
Question 2 .....	5
Answer to Question 2.....	5
Question 3 .....	9
Answer to Question 3.....	9
Question 4 .....	11
Answer to Question 4.....	11
Question 5 .....	12
Answer to Question 5.....	12
Question 6 .....	14
Answer to Question 6.....	14
Question 7 .....	16
Answer to Question 7.....	16
Question 8 .....	20
Answer to Question 8.....	20
Question 9 .....	24
Answer to Question 9.....	24

# Data Structures and Algorithms Questions

## Question 1

Q.12	Which one of the following sequences when stored in an array at locations $A[1], \dots, A[10]$ forms a max-heap?
(A)	23, 17, 10, 6, 13, 14, 1, 5, 7, 12
(B)	23, 17, 14, 7, 13, 10, 1, 5, 6, 12
(C)	23, 17, 14, 6, 13, 10, 1, 5, 7, 15
(D)	23, 14, 17, 1, 10, 13, 16, 12, 7, 5

## Answer to Question 1

### What is a Max-Heap?

A **max-heap** is a special type of **complete binary tree** where:

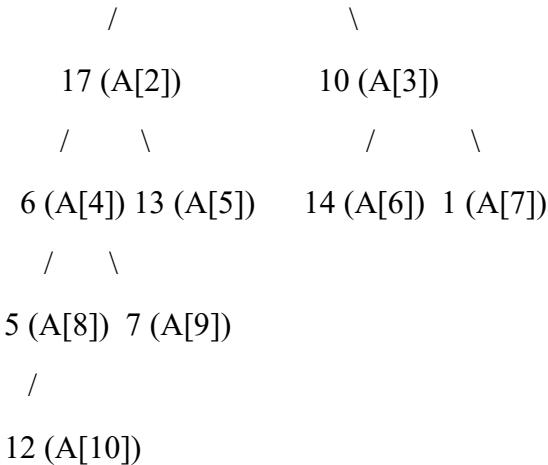
1. **The value at each node is greater than or equal to the values of its children.**
2. **The tree is complete**, meaning all levels are fully filled except possibly the last level, which is filled from left to right.

If the max-heap is stored as an array starting from index **1**, then:

- For any node at index  $i$ :
  - Left child is at index  $2*i$
  - Right child is at index  $2*i + 1$
  - Parent is at index  $i//2$

Lets Check Option A : 23, 17, 10, 6, 13, 14, 1, 5, 7, 12

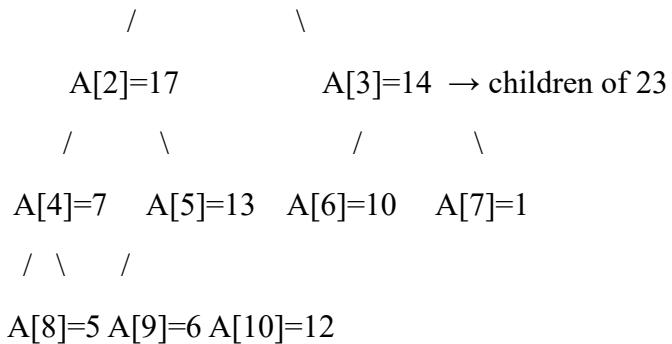
23 (A[1])



Heap property violated here because **14 > 10**, and 14 is a child of 10 . So **Option (A) is NOT a max-heap.**

Option (B): 23, 17, 14, 7, 13, 10, 1, 5, 6, 12

A[1] = 23 → Root node



Example of parent and child nodes relationship :

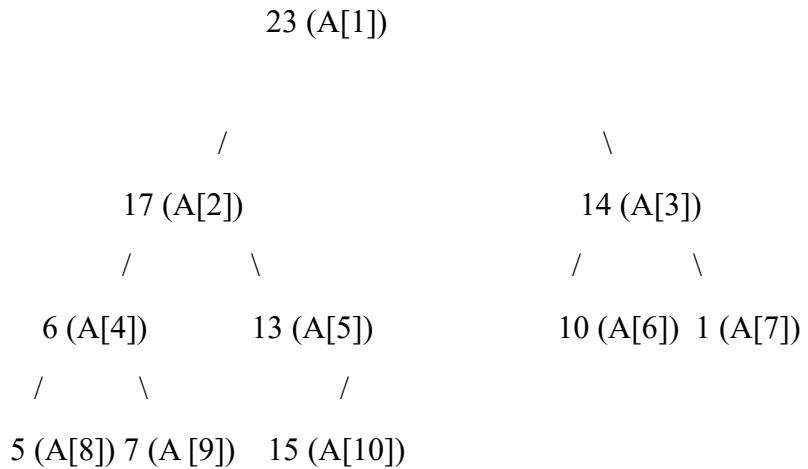
- Each **index** represents a **node**
- The **value at that index** is the **value of the node**
- Parent-child relationship is based on array indices:
- For **node at index i**:
  - **Left child** = node at index  $2*i$
  - **Right child** = node at index  $2*i + 1$

**For instance :**

- Node at index 2 → value = 17
- Left child: index 4 → value = 7
- Right child: index 5 → value = 13

All parent nodes are  $\geq$  their children . This is a **valid max-heap!**

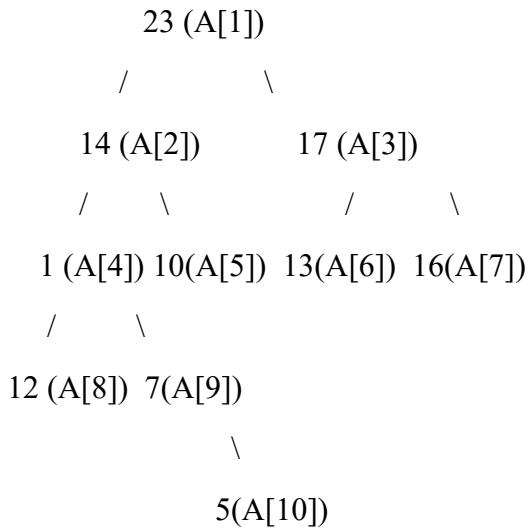
Option (C) : 23, 17, 14, 6, 13, 10, 1, 5, 7, 15



1. **Node A[4] = 6 has child A[9] = 7 → violation of max heap**
2. **Node A[5] = 13 has child A[10] = 15 → violation of max heap**

**multiple violations** exist in Option (C), making it **clearly not a max-heap.**

Option (D): 23, 14, 17, 1, 10, 13, 16, 12, 7, 5



$1 \geq 12, 7$  FALSE . **Violation here:  $1 < 12$  and  $1 < 7$**

So again: **Option (D) is NOT a max-heap .**

## Question 2

Q.13	<p>Let <code>SLLdel</code> be a function that deletes a node in a singly-linked list given a pointer to the node and a pointer to the head of the list. Similarly, let <code>DLLdel</code> be another function that deletes a node in a doubly-linked list given a pointer to the node and a pointer to the head of the list.</p> <p>Let <math>n</math> denote the number of nodes in each of the linked lists. Which one of the following choices is TRUE about the worst-case time complexity of <code>SLLdel</code> and <code>DLLdel</code>?</p>
(A)	<code>SLLdel</code> is $O(1)$ and <code>DLLdel</code> is $O(n)$
(B)	Both <code>SLLdel</code> and <code>DLLdel</code> are $O(\log(n))$
(C)	Both <code>SLLdel</code> and <code>DLLdel</code> are $O(1)$
(D)	<code>SLLdel</code> is $O(n)$ and <code>DLLdel</code> is $O(1)$

## Answer to Question 2

The question is about analysing the **worst-case time complexity** of deleting a node from:

- A singly linked list (**SLL**) – `SLLdel`

- A **doubly linked list (DLL)** – DLLdel

A **singly linked list** is a linear data structure where each element (called a **node**) points to the **next node** in the sequence.

**Structure of a Node:** [ Data | Next ]

- **Data:** stores the value
- **Next:** stores a pointer/reference to the **next node**

**For Example:** [10|] → [20|] → [30|] → null

Each node links to the next one, but **not backwards**.

To reach a node, you must start from the **head** and follow the next pointers.

**Limitations:**

- No direct way to go backward
- To delete a node, you **need access to the previous node**

Lets say for example : [10] → [20] → [30] → [40] → null

You want to delete [30], and you're **only given**:

- A pointer to the **head** of the list
- A pointer to the **node to delete**: [30]

**Why can't you just delete [30]?**

To delete a node in a singly linked list, you need to **update the .next of the previous node**.

In this case, you need to set:

20.next = 40

But you don't have a direct way to access the previous node ([20]) from [30], because SLL only has .next.

**So What Do You Have to Do?**

You must:

1. Start from the **head** node
2. Traverse the list
3. Keep checking each node's .next
4. Stop when you find a node where: current.next == nodeToDelete

Once you find that node (current = 20), then you can do:

current.next = current.next.next;

## What makes it $O(n)$ ?

Let's say you want to delete the node with value  $x$ , and the list has  $n$  elements:

$[1] \rightarrow [2] \rightarrow [3] \rightarrow \dots \rightarrow [x] \rightarrow \dots$

- Worst case:  $x$  is at the **end**
- You must check each node's `.next` to find the one before  $x$
- You might have to scan **all  $n$  nodes**

So it's a **linear scan**, and that makes it  **$O(n)$**

---

## Why Not $O(1)$ ?

$O(1)$  means **constant time** — the operation takes the same time no matter how many elements there are.

To get  $O(1)$  deletion, you'd need:

- Direct access to the **previous node**
- Or the ability to **jump straight to the node before the one to delete**

In SLL, you **can't go backward**, and there are **no shortcuts**, so you **can't do it in  $O(1)$**  unless that previous node is already provided.

---

## Why Not $O(\log n)$ ?

$O(\log n)$  usually happens with **binary search** or **tree-based structures**, where you can eliminate half the data at each step — like:

- Binary Search Tree (BST)
- Balanced trees (e.g., AVL, Red-Black Trees)
- Binary search on arrays

But in an SLL:

- You **can't jump** to the middle
- You **can't divide the list**
- You must go **one node at a time**

So you can't get that logarithmic reduction — it's strictly linear traversal.

## Doubly Linked List (DLL)

A **doubly linked list** is like a singly linked list, but each node has **two** pointers:

- One to the **next** node
- One to the **previous** node

### Structure of a Node:

[ Prev | Data | Next ]

**For Example:** null  $\leftarrow$  [10]  $\leftrightarrow$  [20]  $\leftrightarrow$  [30]  $\rightarrow$  null

Now you can traverse **both forward and backward** through the list.

### Advantages:

- Can go backwards easily
- Can delete a node directly if you have a pointer to it (since you can access both neighbors)

Why is the time complexity of DLLdel (deleting a node from a Doubly Linked List) = O(1)?

Let's break it down step by step

---

### Doubly Linked List (DLL) — Recap

Each node has 3 parts: [ prev | data | next ]

So:

- You can go **forward** (node.next)
  - And **backward** (node.prev)
- 

### Goal: Delete a Node (say, B)

Example DLL: null  $\longleftrightarrow$  [A]  $\longleftrightarrow$  [B]  $\longleftrightarrow$  [C]  $\longleftrightarrow$  null

Each connection:

- A.next = B, B.prev = A
  - B.next = C, C.prev = B
- 

### You Are Given a Pointer to Node B

To delete B, you need to:

B.prev.next = B.next; // A.next = C

B.next.prev = B.prev; // C.prev = A

Now B is completely skipped.

---

### Why Is This O(1)?

Because you already have **direct access** to B and its neighbors (B.prev and B.next) via pointers.

You don't need to:

- Traverse from head
- Search for the node
- Count anything

All pointer updates are done in **constant time**, regardless of list size.

So the answer is Option (D).

## Question 3

Q.20	<p>An algorithm has to store several keys generated by an adversary in a hash table. The adversary is malicious who tries to maximize the number of collisions. Let <math>k</math> be the number of keys, <math>m</math> be the number of slots in the hash table, and <math>k &gt; m</math>.</p> <p>Which one of the following is the best hashing strategy to counteract the adversary?</p>
(A)	Division method, i.e., use the hash function $h(k) = k \bmod m$ .
(B)	Multiplication method, i.e., use the hash function $h(k) = \lfloor m(kA - \lfloor kA \rfloor) \rfloor$ , where $A$ is a carefully chosen constant.
(C)	Universal hashing method.
(D)	If $k$ is a prime number, use Division method. Otherwise, use Multiplication method.

## Answer to Question 3

### Key Terms:

- **Hash table:** A data structure that stores data using a key → hash function → index.
- **Collision:** When two keys end up at the same position in the hash table.
- **Adversary:** A malicious actor who tries to choose keys that cause many collisions.

### **Problem Setup:**

- The attacker can **choose the keys**.
- $k$  = number of keys (greater than number of slots  $m$ ).
- We need a hashing method that makes it **hard for the attacker to force collisions**.

### **What are the options?**

- **(A) Division method:**  $h(k) = k \bmod m$ 
  - Simple and fast.
  - Easy for attacker to guess collisions if  $m$  is known.
- **(B) Multiplication method:**  $h(k) = \text{floor}(m * (kA \bmod 1))$ 
  - More complex, less predictable.
  - Harder for attacker to predict collisions, if  $A$  is well-chosen.
- **(C) Universal Hashing:** Randomized hashing method.
  - **Best defense** against attackers.
  - Uses a randomly chosen function from a family of hash functions.
  - Since the attacker doesn't know the function, it's **very hard to create collisions**.
- **(D) Mix of division and multiplication depending on whether  $k$  is prime.**
  - Irrelevant to the security from attacker; not robust.

### **Correct Answer: (C) Universal hashing method**

#### **Why?**

- Since the adversary **knows** the hash function (or can guess it), methods like division/multiplication are **predictable**.
- **Universal hashing** introduces **randomness** — the hash function is selected at random from a family of functions.
- Because the adversary doesn't know **which** hash function will be used, they **cannot force collisions effectively**.

#### **In Simple Words:**

You want to store keys in a way that an attacker **can't easily make them crash into each other** (collide).

The best way is to pick your hash function **randomly from a secure pool** — so the attacker can't predict how to cause damage.

That's why **Universal Hashing** is the best answer here.

## Question 4

Q.35	The integer value printed by the ANSI-C program given below is _____.
	<pre>#include&lt;stdio.h&gt;  int funcp(){     static int x = 1;     x++;     return x; }  int main(){     int x,y;     x = funcp();     y = funcp() + x;     printf("%d\n", (x+y));     return 0; }</pre>

## Answer to Question 4

### Step-by-step Execution:

#### 1. First Call: `x = funcp();`

- o `static int x = 1;` → initializes x to 1 **only once**.
- o `x++` → x becomes 2.
- o returns 2 → So,  $x = 2$

#### 2. Second Call: `y = funcp() + x;`

- o `funcp()` again:
  - x is already 2 (static retains its value)
  - `x++` → becomes 3
  - returns 3
- o  $y = 3 + x \rightarrow y = 3 + 2 = 5$

#### 3. Final Output:

- o  $x + y = 2 + 5 = 7$

**Answer: 7**

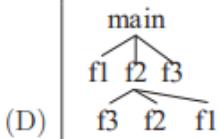
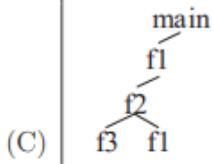
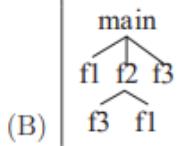
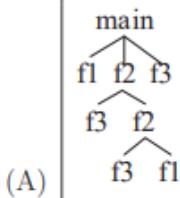
## Question 5

Q.36

Consider the following program:

```
int main()    int f1()    int f2(int X)
{           {           {
    f1();       return(1);   f3();
    f2(2);     }           if (X==1)
    f3();       }           return f1();
    return(0);   }           else
}           }           return (X*f2(X-1));
```

Which one of the following options represents the activation tree corresponding to the main function?



## Answer to Question 5

### What is an Activation Tree?

An **activation tree** (also known as a **function call tree**) is a diagram that shows **how functions are called** during the execution of a program. Each node in the tree represents a function call. A function A that calls function B will have B as its child node in the tree.

It helps visualize **the order and nesting of function calls**, especially when functions call other functions recursively.

You are asked:

"Which option represents the **activation tree** of this program — i.e., which diagram shows how the function calls happen during execution?"

### Let's Simulate What Happens

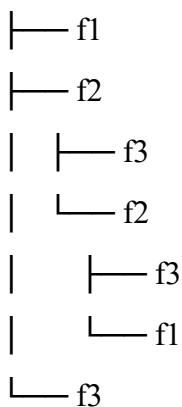
Here's what happens step-by-step in main():

1. f1() is called (simple, no further calls).
2. f2() is called:
  - o It calls f3() first.
  - o Then X != 1, so it calls f2(1):
    - f3() is called again.
    - X == 1, so it calls f1() again.
3. f3() is called again at the end.

If you look at option (A) :

It shows:

main



This perfectly matches the flow we just described! Final Answer: **Option (A)**

## Question 6

Q.46	<p>Let <math>A</math> be a priority queue for maintaining a set of elements. Suppose <math>A</math> is implemented using a max-heap data structure. The operation <math>\text{EXTRACT-MAX}(A)</math> extracts and deletes the maximum element from <math>A</math>. The operation <math>\text{INSERT}(A, \text{key})</math> inserts a new element <math>\text{key}</math> in <math>A</math>. The properties of a max-heap are preserved at the end of each of these operations.</p> <p>When <math>A</math> contains <math>n</math> elements, which one of the following statements about the worst case running time of these two operations is TRUE?</p>
(A)	Both $\text{EXTRACT-MAX}(A)$ and $\text{INSERT}(A, \text{key})$ run in $O(1)$ .
(B)	Both $\text{EXTRACT-MAX}(A)$ and $\text{INSERT}(A, \text{key})$ run in $O(\log(n))$ .
(C)	$\text{EXTRACT-MAX}(A)$ runs in $O(1)$ whereas $\text{INSERT}(A, \text{key})$ runs in $O(n)$ .
(D)	$\text{EXTRACT-MAX}(A)$ runs in $O(1)$ whereas $\text{INSERT}(A, \text{key})$ runs in $O(\log(n))$ .

## Answer to Question 6

### Question Summary:

You're given:

- A **priority queue A** implemented using a **max-heap**
- Two operations:
  - Extract-Max(A): removes and returns the **maximum element**
  - Insert(A, key): inserts a **new element**
- The question is asking about the **worst-case time complexity** of these two operations when A contains n elements.

### Key Concepts: Max-Heap Operations

A **max-heap** is a binary tree where:

- Each parent is  $\geq$  its children
- It's a **complete binary tree** stored as an array

#### 1. Extract-Max(A):

- Removes the root (maximum element).
- To maintain the heap structure:
  - Replace root with last element

- Then "heapify down" (i.e., swap down with the larger child)
- This may go all the way from root to leaf.

**Worst-case time:**

$O(\log n) \leftarrow$  because height of a binary heap with  $n$  nodes is  $\log n$

---

## 2. Insert( $A$ , key):

- Add new key at the end (maintains completeness)
- Then "heapify up" (bubble it up while parent < child)
- This may go from the bottom to the root.

**Worst-case time:**

Also  $O(\log n)$

**Why Option (B) is correct:** (B) Both Extract-Max( $A$ ) and Insert( $A$ , key) run in  $O(\log n)$ .

This exactly matches what we just discussed — both operations involve traversing the **height** of a binary heap, which is  $\log n$  in a complete binary tree.

**Why Option (A) is incorrect:** (A) Both Extract-Max( $A$ ) and Insert( $A$ , key) run in  $O(1)$ .

- $O(1) =$  constant time
- That would only be true if:
  - We removed the root without reheapifying (which would break the heap)
  - We inserted at the end without restoring heap order

**Heap structure would be violated . Not possible in a valid max-heap**

**Why Option (C) is incorrect:** (C) Extract-Max( $A$ ) runs in  $O(1)$  whereas Insert( $A$ , key) runs in  $O(n)$ .

- Extract-Max is **not  $O(1)$**  → we must heapify-down → takes  **$O(\log n)$**
- Insert is **not  $O(n)$**  → heapify-up only involves  $\log n$  levels, not  $n$

Incorrect on both counts

**Why Option (D) is incorrect:** (D) Extract-Max( $A$ ) runs in  $O(1)$  whereas Insert( $A$ , key) runs in  $O(\log n)$ .

- Insert( $A$ , key) is **correctly  $O(\log n)$**
- But **Extract-Max( $A$ ) is NOT  $O(1)$**  — it's  $O(\log n)$  due to heapify-down

Half-correct ≠ correct answer

## Question 7

Q.47	<p>Consider the C function <code>foo</code> and the binary tree shown.</p> <pre>typedef struct node {     int val;     struct node *left, *right; } node;  int foo(node *p) {     int retval;     if (p == NULL)         return 0;     else {         retval = p-&gt;val + foo(p-&gt;left) + foo(p-&gt;right);         printf("%d ", retval);         return retval;     } }</pre> <p>When <code>foo</code> is called with a pointer to the root node of the given binary tree, what will it print?</p> <pre>graph TD; 10((10)) --&gt; 5((5)); 10((10)) --&gt; 11((11)); 5((5)) --&gt; 3((3)); 5((5)) --&gt; 8((8)); 11((11)) --&gt; 13((13));</pre> <p>(A) 3 8 5 13 11 10 (B) 3 5 8 10 11 13 (C) 3 8 16 13 24 50 (D) 3 16 8 50 24 13</p>
------	---

## Answer to Question 7

**What is the question asking?**

You are given:

1. A **binary tree** with integer values
2. A **C function named `foo`**
3. You are asked:

If `foo()` is called with a pointer to the **root node** of the binary tree (value = 10), **what will be printed?**

In other words:

- Understand the **order in which nodes are visited**
- Understand **what values get printed and when**

**Let's Understand the C Code :**

```
typedef struct node {  
    int val;  
    struct node *left, *right;  
} node;
```

A basic structure for a binary tree node:

- val is the value stored in the node
- left and right are pointers to the left and right children

**Function: foo(node \*p)**

```
int foo(node *p) {  
    int retval;  
    if (p == NULL)  
        return 0;  
    else {  
        retval = p->val + foo(p->left) + foo(p->right);  
        printf("%d ", retval);  
        return retval;  
    }  
}
```

Let's break this down:

1. **Base Case:** if (p == NULL) return 0;

If the current node is NULL (empty), return 0.

2. **Recursive Step:**

retval = p->val + foo(p->left) + foo(p->right);

This is the key line!

- **It recursively sums:**
  - the value of the **current node**
  - the result of **left subtree**
  - the result of **right subtree**

So it's doing a **post-order traversal** where:

- You process the **left subtree**
- Then the **right subtree**
- Then calculate and print the total for the **current node**

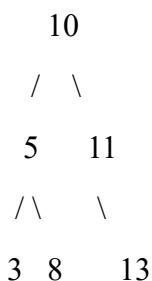
### 3. Print and Return:

```
printf("%d ", retval);
```

```
return retval;
```

- It prints the **sum of the subtree rooted at p**
- Then returns this sum to the parent call

## The Tree Structure



Let's label:

- Root: 10
- Left subtree: 5 → 3, 8
- Right subtree: 11 → 13

## How foo() Executes (Post-order traversal)

Let's simulate the recursion:

### Node 3:

- Left & right are NULL → returns 3
- Prints: 3

### Node 8:

- Left & right are NULL → returns 8
- Prints: 8

### Node 5:

- Left = 3, Right = 8 → total =  $5 + 3 + 8 = 16$
- Prints: 16

### Node 13:

- Left & right NULL → returns 13
- Prints: 13

**Node 11:**

- Left = NULL, Right = 13 → total =  $11 + 0 + 13 = \mathbf{24}$
- Prints: 24

**Node 10 (root):**

- Left = 16, Right = 24 → total =  $10 + 16 + 24 = \mathbf{50}$
- Prints: 50

**Final Output Order:** 3 8 16 13 24 50

This matches **Option (C)**.

**Summary**

**What they're asking:**

“What gets printed when the function foo() is called on the root node?”

**What the code does:**

- **Post-order traversal**
- Calculates the **sum of subtree values**
- Prints the **subtree sum at each node**, bottom-up

**Correct Answer: (C) 3 8 16 13 24 50**

## Question 8

Q.57

Consider the following two-dimensional array D in the C programming language, which is stored in row-major order:

```
int D[128][128];
```

Demand paging is used for allocating memory and each physical page frame holds 512 elements of the array D. The Least Recently Used (LRU) page-replacement policy is used by the operating system. A total of 30 physical page frames are allocated to a process which executes the following code snippet:

```
for (int i = 0; i < 128; i++)
    for (int j = 0; j < 128; j++)
        D[j][i] *= 10;
```

The number of page faults generated during the execution of this code snippet is \_\_\_\_\_.

## Answer to Question 8

### What is demand paging?

- It's a **virtual memory** concept.
- **Pages of a program are not loaded into memory until they're actually used.**
- If the program tries to access something **not currently in memory**, the **operating system raises a "page fault"**.
- Then the OS **loads the required page from disk into RAM**.

### Analogy:

Imagine a huge book. Instead of carrying the entire book, you only carry the chapters you're currently reading. If you need another chapter, you go fetch it.

So, in your program — **only the parts of the array that are actively being used get loaded into RAM**.

### "Each physical page frame holds 512 elements of the array"

### What is a physical page frame?

- It's a **fixed-size block** of physical memory (RAM).
- Think of RAM being divided into equally-sized **slots**.
- Each of these slots is called a **page frame**.

### What does "holds 512 elements" mean?

- The array is made of ints (integers).

- Each frame (slot in memory) can store **512 integer values** from the array.
- So, if you load part of the array into memory, **512 integers (elements) come in at once.**

### For Example:

If your array is  $D[128][128]$ , that's 16,384 integers total.

Each page frame holds 512 integers.

So, to store the entire array, you'd need:

$$16,384 / 512 = 32 \text{ page frames.}$$

But your OS only gives you **30 page frames** to work with. So at any time, **2 pages won't fit**, and **LRU (Least Recently Used) policy** is used to evict old ones.

---

### Putting it all together:

- The array  $D[128][128]$  is **too big to fit entirely in RAM**.
- The OS uses **demand paging** to load only parts of the array **when needed**.
- Each **page fault** happens when a part of the array is accessed but **not currently in memory**.
- And each time, **a block of 512 integers is loaded** into a **page frame**.

### How does RAM work here?

#### RAM = Main Memory

- RAM is **fast memory** where your program runs.
- But it is **limited** in size (e.g., 8 GB, 16 GB).
- So it can't store **everything** at once.

### What Happens:

1. The program starts. The OS puts some **initial pages** into RAM.
2. The program accesses **new data or code**.
3. If it's not already in RAM → **Page fault** occurs.
4. OS takes that missing page from **disk** (virtual memory) → loads it into **RAM**.
5. If RAM is full, the OS may **remove (evict)** an old page from RAM using a policy like LRU (Least Recently Used).

### What is a Page Fault?

A **page fault** occurs when:

The program **tries to access a page** that is **not in RAM**.

**Here's what happens in a page fault:**

1. The CPU tries to access memory (e.g., a variable).
2. The **Memory Management Unit (MMU)** checks the page table.
3. It finds that the required page is **not in RAM**.
4. **Page Fault triggered** — like a red flag!
5. OS **pauses the program**, goes to the **disk**, and **loads that page** into RAM.
6. If RAM is full → OS may **remove an old page** (like closing a tab in a browser).
7. The page is loaded, and the program **resumes exactly where it left off**.

**Analogy:**

Imagine you're reading a book on a tablet with limited memory.

- You start reading Chapter 1 — it's loaded in memory.
- You scroll to Chapter 3 — the tablet goes: "Oh! That's not loaded yet."
- It downloads Chapter 3 from the internet (like fetching a page from disk).
- It might remove Chapter 1 if memory is full.

That "download" = **page fault**.

**Now the actual answer is 4096.**

**What the Official Answer is Assuming:**

They're assuming:

Each page holds 512 elements

Each row has 128 elements

So: 4 rows fit per page (since  $512/128 = 4$ )

Total rows = 128 → so:

$128 / 4 = 32$  pages are needed to store all rows

Now look at the access:

$D[j][i]$  → access 1 element from each row

So in 1 column, we access 128 rows = 32 pages (again)

So: in each of the 128 outer loop iterations, 32 page accesses are made.

Because only 30 pages fit in memory:

In every column, 2 pages must be evicted (as only 30 can stay in RAM)

So every column iteration causes 2 new page faults

But this is not how the answer key got 4096.

**Here's how they did it:**

Each page holds **512 elements**.

Each **access** is to a single element, and there are **16,384 accesses** ( $128 \times 128$ ).

So:

- Each **access** may trigger a page fault, but once a page is loaded, other accesses to it may not.

Because the access pattern jumps across rows ( $D[j][i]$ ), and each page stores **4 rows**, every **i iteration** accesses 128 rows → thus **32 pages**.

So:

- Each **i** (128 total) → accesses **32 pages**
- So: **total page accesses =  $128 \times 32 = 4096$**
- All are **page faults**, because LRU can't retain them all (only 30 pages)

**Answer = 4096 page faults**

**Final Summary:**

- Each column access = 128 elements → 32 different pages
- 128 such columns →  **$128 \times 32 = 4096$**
- That's the total number of page faults when every access triggers a page fault (due to poor locality)
- **Answer = 4096**

## Question 9

Q.59	<p>Consider a sequence <math>a</math> of elements <math>a_0 = 1, a_1 = 5, a_2 = 7, a_3 = 8, a_4 = 9</math>, and <math>a_5 = 2</math>. The following operations are performed on a stack <math>S</math> and a queue <math>Q</math>, both of which are initially empty.</p> <p>I: push the elements of <math>a</math> from <math>a_0</math> to <math>a_5</math> in that order into <math>S</math>.</p> <p>II: enqueue the elements of <math>a</math> from <math>a_0</math> to <math>a_5</math> in that order into <math>Q</math>.</p> <p>III: pop an element from <math>S</math>.</p> <p>IV: dequeue an element from <math>Q</math>.</p> <p>V: pop an element from <math>S</math>.</p> <p>VI: dequeue an element from <math>Q</math>.</p> <p>VII: dequeue an element from <math>Q</math> and push the same element into <math>S</math>.</p> <p>VIII: Repeat operation VII three times.</p> <p>IX: pop an element from <math>S</math>.</p> <p>X: pop an element from <math>S</math>.</p> <p>The top element of <math>S</math> after executing the above operations is _____.</p>
------	---

## Answer to Question 9

You are given:

- A sequence of elements:  
 $a = [1, 5, 7, 8, 9, 2]$   
(These are  $a_0$  to  $a_5$ )
- Two data structures:
  - **Stack S** (Last-In-First-Out)
  - **Queue Q** (First-In-First-Out)
- Initially both S and Q are empty.

You are told to:

- Perform a sequence of operations (I through X)
- After all operations, tell **what is the top element of the stack S**

### Step-by-step execution

#### Step I: Push $a_0$ to $a_5$ into S

Sequence: 1, 5, 7, 8, 9, 2 → push into **Stack S**

Since stack is **LIFO**, top is the last pushed.

After step I:

$S = [1, 5, 7, 8, 9, 2]$  ← 2 is on top

---

### **Step II: Enqueue $a_0$ to $a_5$ into Q**

Sequence: 1, 5, 7, 8, 9, 2 → enqueue into **Queue Q**

Since queue is **FIFO**, front is the first enqueued.

After step II:

$Q = [1, 5, 7, 8, 9, 2]$  ← 1 is at front

---

### **Step III: Pop from S**

$S.pop()$  → removes 2

Now:  $S = [1, 5, 7, 8, 9]$

---

### **Step IV: Dequeue from Q**

$Q.dequeue()$  → removes 1

Now:  $Q = [5, 7, 8, 9, 2]$

---

### **Step V: Pop from S**

$S.pop()$  → removes 9

Now:  $S = [1, 5, 7, 8]$

---

### **Step VI: Dequeue from Q**

$Q.dequeue()$  → removes 5

Now:  $Q = [7, 8, 9, 2]$

---

### **Step VII: Dequeue from Q and Push to S**

Dequeue → 7, Push to S

Now:

- $Q = [8, 9, 2]$
  - $S = [1, 5, 7, 8, 7]$
- 

### **Step VIII: Repeat VII three times**

1. Dequeue 8, Push to S → S = [1, 5, 7, 8, 7, 8], Q = [9, 2]
  2. Dequeue 9, Push to S → S = [1, 5, 7, 8, 7, 8, 9], Q = [2]
  3. Dequeue 2, Push to S → S = [1, 5, 7, 8, 7, 8, 9, 2], Q = []
- 

### **Step IX: Pop from S**

S.pop() → removes 2  
Now: S = [1, 5, 7, 8, 7, 8, 9]

---

### **Step X: Pop from S**

S.pop() → removes 9  
Now: S = [1, 5, 7, 8, 7, 8]

---

**Final Answer:**

**Top element of S = 8**

**Answer: 8**